

# A tutorial on Bayes' theorem application to physical model parameter extraction

Weiyao Ke

modified from JetScape Winter School 2019 talk by Jake Coleman

Dec 3, 2019

1 Bayes' theorem

2 Complex model: high dimensional input and output

- Designing parameter samples
- Interpolator: Gaussian process
- Dimensional reduction: principal component analysis

3 Application: a toy model for jet-quenching

# Bayes' parameter extraction

The problem:

1. Given a model  $\mathcal{M}$ : compute quantities  $\mathbf{y}$  with input parameters  $\mathbf{x}$ .
  2. Given a prior belief of  $\mathbf{x}$ ' true value's distribution  $P_0(\mathbf{x}_{\text{true}})$
  3. Given observations  $\mathbf{y}_{\text{exp}}$ .
- !! Ask for the updated probability distribution of  $\mathbf{x}_{\text{true}}$ :  $P(\mathbf{x}_{\text{true}})$ .

Bayes' theorem:

$$\underbrace{P(\mathbf{x}_{\text{true}}|\mathcal{M}, \mathbf{y}_{\text{exp}})}_{\text{Posterior}} = \frac{\overbrace{P_L(\mathbf{y}_{\text{exp}}|\mathcal{M}, \mathbf{x}_{\text{true}})}^{\text{Likelihood}} \overbrace{P_0(\mathbf{x}_{\text{true}})}^{\text{Prior}}}{\underbrace{\int P_L(\mathbf{x})P_0(\mathbf{x})d\mathbf{x}}_{\text{Normalization (evidence)}}}$$

Often the form  $P_L$  is unknown. Commonly assumed to take the form:

$$\ln P_L = C - \frac{1}{2}\Delta\mathbf{y}\Sigma^{-1}\Delta\mathbf{y}^T, \quad \Delta\mathbf{y} = \mathbf{y}_{\text{exp}} - \mathbf{y}(\mathbf{x}; \mathcal{M}), \Sigma : \text{Contains uncertainties}$$

## A toy model: flip coin

The problem:

- A coin flip experiments gets  $m = 7$  head out of  $n = 10$  trials.
- What is the probability  $\theta$  to get a head each time.

The model:

- Outcome follows a binomial distribution. We know the exact likelihood function  $P_L(m \text{ out of } n | \theta) = C_n^m \theta^m (1 - \theta)^{n-m}$ .
- Try different prior:  $P_0(\theta) \propto \theta^{\alpha-1} (1 - \theta)^{\beta-1}$

The resulting inference on  $p$ :

- Posterior probability distribution  $P(\theta) \propto \theta^{m+\alpha-1} (1 - \theta)^{n-m+\beta-1}$
- Mean  $\pm$  Std:  $\theta = \frac{m+\alpha}{n+\alpha+\beta} \pm \sqrt{\frac{(m+\alpha)(n-m+\beta)}{(n+\alpha+\beta)^2(n+\alpha+\beta+1)}}$  ( $P(\theta)$  has all the information).
- How do  $\alpha, \beta$  change the inference?

# For complex model, it is impossible to evaluate outcome at arbitrary input

Computationally intensive to evaluate the full model.

- For example, 2+1D event-by-event viscous-hydro based HIC simulation  $\sim 5\text{-}10$  events/h. To compute observables with one set of input take  $\mathcal{O}10^4$  minimum biased events  $\rightarrow 1\text{-}2\text{kh}$
- Compute at finite number of input points and interpolate the function  $\mathbf{y}(\mathbf{x}; \mathcal{M})$

Two problems of interpolation:

1. How to efficiently interpolate high-dimensional input  $\mathcal{O}(10)$ ?
2. How to efficiently interpolate a function with vector output.
  - ▶ For example  $\mathbf{y} = (\langle E_T \rangle, dN_{\text{ch}}/d\eta \langle p_T \rangle, v_2, v_3, v_4) \times [8 \text{ centrality}] \times [2 \text{ collision systems}]$  gives an 80-dimensional output at each input.

# Latin hyper-cube sampling: optimized samples of high dimensional parameter space

Assume the model is well-behaved in the parameter space.

- Grid samples require too many samples  $N \propto n^d$
- Random samples result in over populated regions and large gap.

Latin hyper-cube sampling (LHS): random sampling with

- Marginalized distribution is uniform.
- Maximize the minimum distance between two samples.
- Usually  $N \propto d$

# Gaussian emulator: fast interpolator

Given  $y$  values at finite number of parameter sets  $\mathbf{x}_i$ . Gaussian process is:

- A non-parametric interpolation.
- $y^*$  at a new point  $\mathbf{x}^*$  are inferred from its correlation with all other points ( $x_i$ ).
- It also estimates the interpolation uncertainty!

How does it work:

- Assume  $y(\mathbf{x}^*)$ ,  $y(\mathbf{x}_i)$  form a multi-variated Gaussian distribution (centered):

$$\begin{bmatrix} y^* \\ y_i \end{bmatrix} \sim \mathcal{N} \left( \mu = 0, \text{cov} = \begin{bmatrix} k(\mathbf{x}^*, \mathbf{x}^*), k(\mathbf{x}^*, \mathbf{x}_i) \\ k(\mathbf{x}_i, \mathbf{x}^*), k(\mathbf{x}_i, \mathbf{x}_j) \end{bmatrix} \right), \quad k(\cdot, \cdot) \text{ is a kernel function}$$

- $P(y^*)$  is simply the conditional probability  $P(y^*) = \frac{P(y, y_i)}{\int P(y, y_i) dy} \Big|_{y_i=y(\mathbf{x}_i)}$ ,

$$y^* \sim \mathcal{N} \left( \mu = k(\mathbf{x}^*, \mathbf{x}_i) k^{-1}(\mathbf{x}_i, \mathbf{x}_j) y_j, \text{cov} = k(\mathbf{x}^*, \mathbf{x}^*) - k(\mathbf{x}^*, \mathbf{x}_i) k^{-1}(\mathbf{x}_i, \mathbf{x}_j) k(\mathbf{x}_j, \mathbf{x}^*) \right)$$

- Question: what happens when  $\mathbf{x}^* \rightarrow \mathbf{x}_i$

# Gaussian emulator: kernel function and training

One do have freedom in chosen the kernel function and its “hyperparameters”, a common choice is a Gaussian-like correlation function:

- $k(\mathbf{x}_i, \mathbf{x}_j) = \sigma^2 \exp(-\frac{1}{2}|(\mathbf{x}_i - \mathbf{x}_j)/\mathbf{L}|^2)$ : two point correlation  $\langle y(\mathbf{x}_i)y(\mathbf{x}_j) \rangle$
- $\sigma, \mathbf{L}$  called hyperparameters: They are, in principle, unknown. But one can use a set of “optimized values” by balancing the quality of the fit and the complexity of the emulator  
→ training process
- Question: what happens if  $L$  is too small / too large.



# Principal component analysis: data reduction for vector function

Suppose one has computed a list of observables at  $n$  design point,

$$\mathbf{x}_1 \rightarrow \mathbf{y}_1 = [y_1(\mathbf{x}_1), \dots, y_m(\mathbf{x}_1)]$$

...

$$\mathbf{x}_n \rightarrow \mathbf{y}_n = [y_1(\mathbf{x}_n), \dots, y_m(\mathbf{x}_n)]$$

Inefficient to construct  $m$  emulator for each quantity. Data has intrinsic correlation.

- For example, if  $y_i$  and  $y_j$  has a strong linear correlation, then there is effectively only one degree of freedom ( $z_1(\mathbf{x}) = ay_1(\mathbf{x}) + by_2(\mathbf{x}), z_2 = by_1 - ay_2 \ll \mathcal{O}(z_1)$ ) instead of two.
- $(y_1(\mathbf{x}), y_2(\mathbf{x})) \xrightarrow{\text{Change basis}} (z_1(\mathbf{x}), z_2(\mathbf{x})) \xrightarrow{\text{Dim reduction}} (z_1(\mathbf{x}))$
- Principal component analysis (PCA): the statistical tool to generalize this simple example.

# Principal component analysis: data reduction for vector function

- PCA: systematically construct the new basis by diagonalize the empirical covariance matrix

$$C_{ij} = \frac{1}{n} \sum_{k=1}^n \underbrace{\frac{y_i(\mathbf{x}_k) - \bar{y}_i}{\sigma_i}}_{\text{Standardized data } Y_{ik}} \frac{y_j(\mathbf{x}_k) - \bar{y}_j}{\sigma_j} = \frac{1}{n} \mathbf{Y} \mathbf{Y}^T$$

$$\mathbf{C} = \mathbf{V} \text{Diag}\{\lambda_1, \lambda_2, \dots\} \mathbf{V}^T, \quad \lambda_1 > \lambda_2 > \dots$$

- The new basis becomes  $z_i(\mathbf{x}) = V_{ij}^T y_j(\mathbf{x})$ ,  $i = 1, 2, \dots, m$ , with variance  $\lambda_i$ .
- Not all of them are important, one usually truncates at  $N \ll m$  and construct emulators for the first  $N$  most important principal component.
- Given  $z_{i \leq N}$ ,  $y_i$  can be reconstructed  $y_i = V_{ij} z_{j \leq N}$ .

# A toy-model example

## A toy-model of jet-quenching

- Initial spectrum

$$\frac{dN_0}{dp_T} \propto \frac{p_T}{(3^2 + p_T^2)^3}$$

- Energy loss  $\Delta E$  follows a  $\Gamma$ -distribution,

$$P(\Delta E) \propto \Delta E^{\mu^2/\sigma^2-1} e^{-\mu\Delta E/\sigma^2}$$
$$\mu = A\sqrt{p_T}, \sigma = B\mu$$


- The quenched spectrum

$$\frac{dN_1}{dp_T}(p_T) = \int \frac{dN_0}{dp_T}(p_T + x) P(x) dx$$

## A toy-model example


- Assume the model is perfect and the truths are:  $A = 1.0$ ,  $B = 0.5$ .
- A measurement with finite statistics (5%) and systematic bias (10%).

$$y_{\text{exp}} \approx y_{\text{true}} \pm \sigma_{\text{stat}} \pm \sigma_{\text{sys}}$$






## Toy model: make design

An 50-point design with  $A \in [0.05, 3]$  and  $B \in [0.05, 1.5]$ .  
Model calculations of  $R_{AA}$  widely spread between 0 and 1.




# Toy model: PCA

The first two PCs account for more than 99.8% of the data variance.  
PC1 is an overall shift, PC2 and PC3 capture the shapes.




## Toy model: Validation

Compare GPs' predictions to model calculations at new parameter sets. Relative error on the right.




# Toy model: Posterior for parameters

Using uncorrelated sys error



Using correlated sys error




\* Mistreating correlated uncertainty may bias the credible region.




# Toy model: Posterior predictions

Using uncorrelated sys error



Using correlated sys error



\* Mistreating correlated uncertainty may lead to overconfident prediction uncertainties.

## Recap of Whole Analysis

- Pick design points via a Latin Hypercube, run the computer model at those design points.
- Transform the computer output via PCA, pick  $R$  principal components.
- Pick a covariance function, and train  $R$  independent GPs on the first  $R$  columns of the PCA-transformed computer model output.
- Perform calibration, getting posterior draws for input parameters.
  - ▶ For each  $\theta$  draw, find the GP predictions, transform them back from PCA, then put those values in the likelihood.